

Pthreads (POSIX Threads)

Dr. Gen Kawamura
ATLAS Experiment, A. Quadt

II. Physikalisches Institut, Georg-August-Universität Göttingen

Practical Course on Parallel Computing Apr., 2018



Bundesministerium
für Bildung
und Forschung



BMBF-Forschungsschwerpunkt
ATLAS-EXPERIMENT

Physik bei höchsten Energien mit dem ATLAS-Experiment am LHC

FSP 103

ATLAS

Processes and Threads

- Processes

- Threads

POSIX Threads

- General Concepts

- Create, Exit and Cancel Threads

- Shared Data

- Locking Data

- Signaling and Condition Variables

Performance

- Performance Considerations

- Bug and Performance Example

Conclusion



Processes and Threads

- Processes

- Threads

POSIX Threads

- General Concepts

- Create, Exit and Cancel Threads

- Shared Data

- Locking Data

- Signaling and Condition Variables

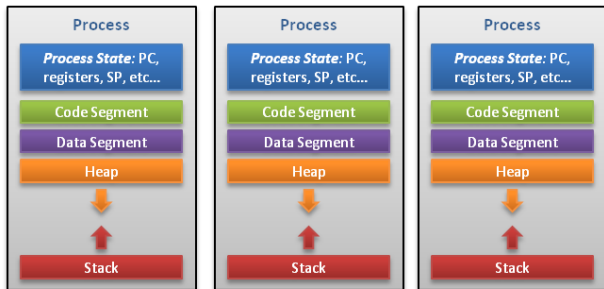
Performance

- Performance Considerations


- Bug and Performance Example

Conclusion


- Common operating systems on common hardware handle a lot of *processes* at once
- Each one with it's own set of *virtual memory*
- All processes are strictly separated for security, simplicity and compatibility reasons
- Each *core* can only run one process at a time
- At different intervals the *scheduler* stops the process, changes the memory-Content of the CPU core (aka *the registers*) and starts another one.



Memory structure of a process on a common Operating System/ CPU architecture



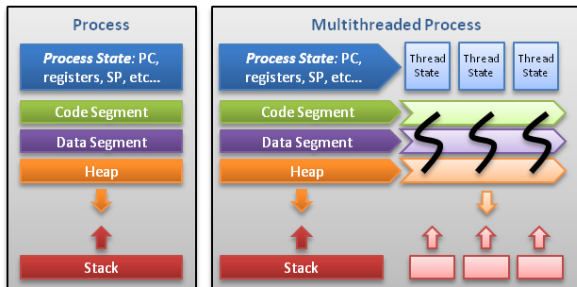
a) shared memory



b) distributed memory

- Run one program per CPU core simultaneously
 - Start one process which sets up the environment
 - and spawns one worker-process per core via `fork(3)`
- Idea: Communicate via inter-process (shared) memory
- But ...
 - Memory from different processes strictly separated (on common OS)
 - How to deal with simultaneous access on same memory page?

- Idea for SMP machines: separate the state and stack but share the heap
- *Lightweight Process or Thread*



© Alfred Park, <http://randu.org/tutorials/threads>

Threads contain only necessary information, such as a stack (for local variables, function arguments, return values), a copy of the registers, program counter and any thread-specific data to allow them to be scheduled individually. Other data is shared within the process between all threads.

- How does the OS scheduler handle threads?
 - 1:1 i.e. one thread is one *job* in the scheduler. The case for most recent OS
 - n:1 or n:m i.e. multiple threads are mapped to one *job*. In this case the library or even the threads have to schedule themselves
- One might compare the 2 models above with preemptive vs. cooperative scheduling or with kernel- vs. user-(space)-threads.

- Linux: clone (2) system call; implemented in the NPTL-lib (and glibc)
- Mac OS X: NSThread class (from Cocoa)
- Windows: CreateThread() library call
- POSIX Threads: *pthread*
 - most Unix': Linux, Mac OS X, Solaris, BSDs...
 - even on Windows
- And many abstractions like boost, QT, glib etc.

We will use POSIX Threads (pthreads)



- Set of (c-) library functions in pthread.h
- Abstracts the underlying OS
- Provides very basic functionality but everything needed to start
- If using frameworks like QT one should use their implementations.
- For C++ one can use the language inherent std::thread class (since C++11)
- C version 11 also has standard threads.h but this is not widely implemented

- ...we will have a close look at some of pthreads features
- ...we will learn about general concepts of multi-threading

Processes and Threads

Processes

Threads

POSIX Threads

General Concepts

Create, Exit and Cancel Threads

Shared Data

Locking Data

Signaling and Condition Variables

Performance

Performance Considerations

Bug and Performance Example

Conclusion

- Not all features are available on all systems
- You only deal with functions
- If you want to change data objects you have to (!) use special functions

Working with pthreads

```
#include <pthread.h>
pthread_t thread;
// Create attribute object
pthread_attr_t attr;
// Initialize it
pthread_attr_init(&attr);
// Change it
pthread_attr_setdetachstate(&attr, PTHREAD_CREATE_JOINABLE);
// use it
r = pthread_create(&thread, &attr, [...]);
```

Basis

```
//compile with gcc -pthread basic.c
#include <stdio.h>
#include <pthread.h>
void *hello()
{
    printf("Hello World.\n");
    pthread_exit(NULL);
}
main () {
    pthread_t thread;
    pthread_create(&thread, NULL, hello, NULL);
    pthread_exit(0);
}
```

- A thread is terminated if
 - the function ends
 - it calls *pthread_exit*(int return_value)
 - it gets killed with *pthread_cancel*(thread_id)
 - main() ends without waiting (it might wait with *pthread_exit*)
 - the process is terminated/ killed by the OS
- *pthread_exit* does not clean after itself – you have to free() memory, close files etc.

- On Linux you can change the scheduling parameters via `setpriority` (2), `pthread_setaffinity_np` (3) or `sched_setaffinity` (2).
- This might be important for binding on a specific core on NUMA machines.

With arguments


```
void *answer(void *value)
{
    long number = (long) value;
    printf("The answer is %ld.\n", number);
    pthread_exit(NULL);
}

int main () {
    long value = 42;
    pthread_t thread;
    pthread_create(&thread, NULL, answer, (void *) value);
    pthread_exit(0);
}
```

With arguments

Read `hello_arg2.c`

- `pthread_join(threadid, status)` – waits for thread *threadid* and writes the `pthread_exit` return code into *status*
- A thread can only be joined by exactly one other thread
- Example later!



- Global C variables are global over thread boundaries
- Memory in heap (*malloc*) is global over threads boundaries
- Variables in stack are *non-global*

Shared Data with a pointer

```
int global_variable=42;
int main() {
    int non_global_variable=23;
    int *pointer_to_global_data = malloc(sizeof(int));
    [...]
}
```

Basic Example

```
int answer=42;
void *hello ()
{
    printf("The answer is again %d\n", answer);
}
int main () {
    pthread_t thread;
    pthread_create(&thread, NULL, hello, NULL);
    pthread_create(&thread, NULL, hello, NULL);
    pthread_exit(0);
}
```

- pthread offers a native Mutex implementation
- A Mutex shields a part of code. Once you are in this code no one else can go into it.
- Syntactically a Mutex is a set of functions and a data object

```
mutex_t mymutex;  
void thread_1() {  
    [...]  
    lock(mymutex);           // as long as mymutex is locked  
    do_something();         // no one else can lock it  
    unlock(mymutex);  
    [...]  
}  
void thread_2() {  
    [...]  
    lock(mymutex);           // so thread_2 might have to wait  
    do_something_else();  
    unlock(mymutex);  
    [...] }
```

- *pthread_mutex_t* – mutex data structure
- *pthread_mutex_init*(mutex, NULL) – initialize mutex variable
- *pthread_mutex_destroy*(mutex) – destroy it
- *pthread_mutex_lock*(mutex) – lock the mutex; will block and stop the thread until mutex is available
- *pthread_mutex_unlock*(mutex) – can only be called by the mutex-owning thread
- *pthread_mutex_trylock*(mutex) – lock the mutex but does not block; might return with a *impossible-to-lock* error code

Example for the use of a mutex



GEORG-AUGUST-UNIVERSITÄT
GÖTTINGEN

dotprod_serial.c
dotprod_mutex.c

- A mutex does no magic! The one and only function is to block until the owner unlocks it.
- Think of it as a gentleman's agreement.
- The scheduling is non-deterministic! Any thread might get the lock first! Beware of deadlocks!!

- a semaphore is similar to a mutex but it *counts* the number of lock holders
- pthread does not offer a native semaphore type!
- but *semaphore.h* does.
- c.f. `sem_init` (3)

- With a condition variable we can signal another thread about an event, e.g. that we are finished doing something.

```
physics() {  
    calculate_gravity();  
    signal_ready(physics);  
    [...] }  
artificial_intelligence() {  
    look_around();  
    move_oponents();  
    signal_ready(ai);  
    [...] }  
game(){  
    for_each_timestep {  
        [...]  
        wait_for(physics);  
        wait_for(ai);  
        paint_graphics();  
    }  
}
```

- *pthread_cond_t* – data structure
- *pthread_cond_init*(condition, NULL)
- *pthread_cond_destroy*(condition)
- *pthread_cond_wait*(condition, mutex) – wait for condition *condition*
- *pthread_cond_signal*(condition) – signal to one thread only that the *condition* is fulfilled
- *pthread_cond_broadcast*(condition) – signal to EVERYBODY that the *condition* is fulfilled

- You always need an additional mutex to shield the condition!
- `cond_wait(cond, mutex)`
 - should be called after *mutex* is locked by the same thread!
 - *unlocks mutex* and blocks
 - waits for the *cond* to be signaled
 - unblocks and immediately *locks mutex*
 - You have to unlock the *mutex* afterwards!
- The thread might wake up from `cond_wait` although the condition is not fulfilled! → You should put it inside a *while*-loop.
- If there is the smallest possibility that more then one thread waits for a condition – use `cond_broadcast`!

condvar.c

- A barrier is a construct to synchronize threads
- All threads that arrive at a barrier have to wait until everybody else is there
- When initializing we have to specify the maximum number of threads that have to wait

```
mybarrier = barrier(9); // we have 8 planets and one sun

calc_planet_position (my_planet) {
    while (true) {
        calc_force_on_my_planet(all_planet_positions);
        move_my_planet();

        //wait for the other planets to finish
        barrier_wait(mybarrier);
    }
}
```


- `pthread_barrier_t` – data type
- `pthread_barrier_init`(barrier, attr, number) – initialize new barrier which stops *number* of threads
- `pthread_barrier_wait`(barrier) – blocks until *number* threads called this function
- `pthread_barrier_destroy`(barrier)

Processes and Threads

Processes

Threads

POSIX Threads

General Concepts

Create, Exit and Cancel Threads

Shared Data

Locking Data

Signaling and Condition Variables

Performance

Performance Considerations

Bug and Performance Example

Conclusion

- *Lock granularity* – How coarse or fine are your mutexes? Do they lock a whole structure or fields of a structure? The more fine-grained, the more concurrency you can gain.
- *Lock frequency* - Are you locking (too) often? Locking at unnecessary times? Reduce such occurrences to fully exploit concurrency and reduce synchronization overhead.
- *Critical sections* - You should minimize critical sections i.e. section that can only be entered by one thread at a time.
- *Worker thread pool* - If you are using a Boss/Worker thread model, make sure you pre-allocate your threads instead of creating threads on demand.
- *Too many threads?* - At what point are there too many threads? Can it severely impact and degrade performance?

bug6.c
bug6_correct.c

Thread 1

```
pthread_mutex_lock(&m1);  
/* use resource 1 */  
  
pthread_mutex_lock(&m2);  
  
/* use resources 1 and 2 */  
  
pthread_mutex_unlock(&m2);  
pthread_mutex_unlock(&m1);
```

Thread 2

```
pthread_mutex_lock(&m2);  
/* use resource 2 */  
  
pthread_mutex_lock(&m1);  
  
/* use resources 1 and 2 */  
  
pthread_mutex_unlock(&m1);  
pthread_mutex_unlock(&m2);
```

Processes and Threads

Processes

Threads

POSIX Threads

General Concepts

Create, Exit and Cancel Threads

Shared Data

Locking Data

Signaling and Condition Variables

Performance

Performance Considerations

Bug and Performance Example

Conclusion

- Threads are a nice way to parallelize problems on a shared memory architecture
- pthreads offer an OS abstraction for threads, mutexes and signal handling (conditions and barriers)
- Avoid race conditions but also avoid deadlocks.

- pthreads (7) man page and pthread_* (3) man pages
- http://pages.cs.wisc.edu/~travitch/pthreads_primer.html
- <http://randu.org/tutorials/threads/>
- <https://computing.llnl.gov/tutorials/pthreads/> (also has a reference)
- Practical Course on Parallel Computing - Sose2015 (The most content/codes were from this material)

